


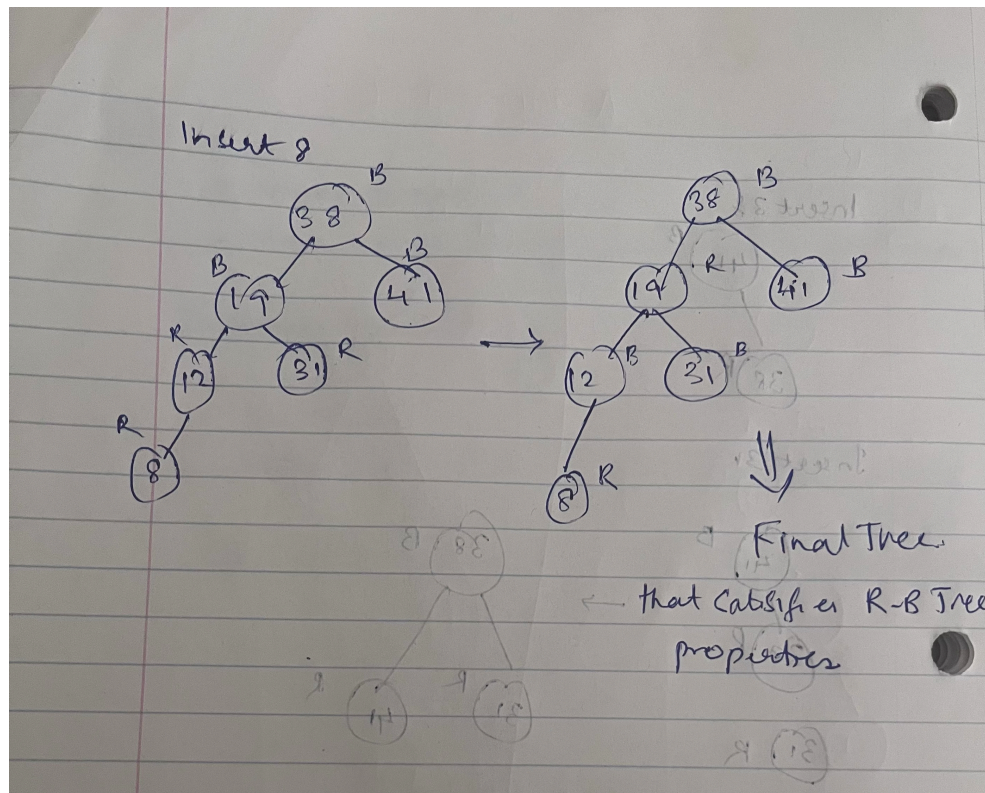
ASSIGNMENT 9

Vikramaditya Reddy

Z1973679

1. Show the red-black trees that result after successively inserting each of the keys 41, 38, 31, 12, 19, 8 into an initially empty red-black tree. For each key, show each tree after inserting each key.





2. Argue that any arbitrary n -node binary search tree can be transformed into any other arbitrary n -node binary search tree using $O(n)$ rotations. Hint: First show that at most $n-1$ right rotations suffice to transform the tree into a right-going chain.)

Consider any arbitrary binary search tree with n nodes. We can transform it into a right-going chain in the following way:

- Start at the root node and move down the left child as far as possible until we reach a leaf node.
- Perform a right rotation to bring the leaf node up to the root of the right subtree.
- Repeat above steps until the tree becomes a right-going chain.

Each iteration reduces the height of the left subtree by one and finally we end up with a right-going chain with n nodes. Since the original tree has a height of at most $n-1$, we need at most $n-1$ iterations and therefore at most $n-1$ right rotations to transform it into a right-going chain.

Now, let's consider two right-going chains, A and B, each with n nodes. To transform A into B, we can use the following algorithm:

- Start at the root of A and B.
- If the roots have the same key, move to their right children.
- If the root of A has a smaller key than the root of B, perform a left rotation on the root of B and its right child. Then move to the left child of the new root of B.
- If the root of A has a larger key than the root of B, perform a right rotation on the root of A and its left child. Then move to the right child of the new root of A.
- Repeat until the roots of A and B have the same key.

Here we can perform each rotation in $O(1)$ time.

Therefore, the total number of rotations required to transform A into B is at most $2(n-1)$, which is $O(n)$.

3. Can the black-heights of nodes in a red-black tree be maintained as a data field in the nodes of the tree without affecting the asymptotic performance of insertion and search in the red-black tree? Show how, or argue why not.

Ans:

Yes, the black-heights of nodes in a red-black tree can be maintained as a data field in the nodes of the tree without affecting the asymptotic performance of insertion and search in the red-black tree.

Red-black trees are a type of self-balancing binary search tree that ensure that the tree is balanced and that the worst-case performance of operations is $O(\log n)$.

The black-height of any node in the tree is the same for all nodes in its subtree. The black-height of a node is defined as the number of black nodes on the path from that node to a leaf node, not including the node itself.

Maintaining the black-heights of nodes as a data field in the nodes of the tree does not affect the asymptotic performance of insertion and search in the red-black tree. This is because the black-heights of nodes can be updated during insertion or deletion operations in $O(1)$ time, which does not affect the $O(\log n)$ worst-case performance of these operations.

4.The interval tree is also a red-black tree. When inserting or deleting a node from interval tree, rotation will be involved. The `max` data fields of involved nodes will need update. Use Left- Rotate to explain how the `max` data fields of involved nodes can be updated in $O(1)$ time.

Left- Rotate method to update max data fields involved nodes in $O(1)$ time

- Let a be the node that will be rotated to the left, and b be its right child.
- Save the max data fields of a, b, and b's left child (if it exists).
- Perform the left rotation as usual, with a becoming b's left child and b's left child becoming a's right child.
- Update the max data fields of a and b based on their new children.
- If b has a left child, update its max data field based on its new children.
- Restore the saved max data fields from step 2 to the appropriate nodes.

5.Describe an efficient algorithm that, given an interval i , returns an interval overlapping i that has the minimum low endpoint, or null if no such interval exists.

- Create an empty variable to store the minimum low endpoint interval.
- Iterate through each interval in the dataset
- if overlaps with i , check if variable is empty or if the low endpoint of the current interval is lower than the low endpoint of the result interval.
- If either condition in previous step is true, update the result variable to be the current interval.
- Once all intervals have been checked, return the result variable if it is not empty, otherwise return null.

Algorithm:

function Overlap(i , intervals):

$x = \text{null}$

 for interval in intervals:

 if $i.\text{high} \geq \text{interval}.\text{low}$ and $\text{interval}.\text{high} \geq i.\text{low}$:

 if x is null or $\text{interval}.\text{low} < x.\text{low}$:

$x = \text{interval}$

 return x

